

Entrée [1]:

```
# Badr TADJER / Naoufal ARRADI / Leo TRAN  
# M1-LS1-APP
```

Predict survival on the Titanic

In this Lab, we ask you to apply the tools of machine learning to predict which passengers survived the tragedy

Dataset

The dataset contains 891 observations of 12 variables:

- **PassengerId**: Unique ID for each passenger
- **Survived**: Survival (0 = No; 1 = Yes)
- **Pclass**: Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
- **Name**: Name
- **Sex**: Sex
- **Age**: Age
- **Sibsp**: Number of Siblings/Spouses Aboard
- **Parch**: Number of Parents/Children Aboard
- **Ticket**: Ticket Number
- **Fare**: Passenger Fare
- **Cabin**: Cabin
- **Embarked** Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)

Entrée [2]:

```
# imports  
import warnings  
warnings.filterwarnings('ignore')  
import pandas as pd  
import numpy as np  
# your code here
```

Entrée [3]:

```
titanic = pd.read_csv('titanic.csv', index_col = 0)
titanic.head()
```

Out[3]:

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
PassengerId										
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85
3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123
5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN

Entrée [4]:

```
# print some info about the dataframe
# your code here
titanic.info
```

Out[4]:

```
<bound method DataFrame.info of
PassengerId      Survived  Pclass \
1                0         3
2                1         1
3                1         3
4                1         1
5                0         3
...            ...      ...
887              0         2
888              1         1
889              0         3
890              1         1
891              0         3

Age \
PassengerId      Name      Sex
1      Braund, Mr. Owen Harris  male
22.0
2      Cumings, Mrs. John Bradley (Florence Briggs Th...  female
38.0
3      Heikkinen, Miss. Laina  female
26.0
4      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female
35.0
5      Allen, Mr. William Henry  male
35.0
...            ...      ...
...            ...      ...
887      Montvila, Rev. Juozas  male
27.0
888      Graham, Miss. Margaret Edith  female
19.0
889      Johnston, Miss. Catherine Helen "Carrie"  female
NaN
890      Behr, Mr. Karl Howell  male
26.0
891      Dooley, Mr. Patrick  male
32.0

SibSp  Parch      Ticket      Fare Cabin Embarked
PassengerId
1      1      0      A/5 21171      7.2500  NaN      S
2      1      0      PC 17599     71.2833  C85      C
3      0      0  STON/O2. 3101282      7.9250  NaN      S
4      1      0      113803     53.1000  C123     S
5      0      0      373450      8.0500  NaN      S
...      ...      ...      ...      ...      ...
887      0      0      211536     13.0000  NaN      S
888      0      0      112053     30.0000  B42      S
889      1      2      W./C. 6607     23.4500  NaN      S
890      0      0      111369     30.0000  C148      C
891      0      0      370376      7.7500  NaN      Q
```

```
[891 rows x 11 columns]>
```

Looks like there are some Nan values, let's see how many for each column

Entrée [5]:

```
titanic.isnull().sum()
```

Out[5]:

```
Survived      0
Pclass        0
Name           0
Sex            0
Age          177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin        687
Embarked       2
dtype: int64
```

Cabin contains a lot of Nan values, we'll drop this column

We'll replace the Nan values in **Age** with the age's median, and the ones in **Embarked** with 'S', which is the most frequent one in this column

Entrée [6]:

```
# your code here to drop Cabin
titanic = titanic.drop(columns = "Cabin")
# check the fillna documentation: http://pandas.pydata.org/pandas-docs/stable/genera
titanic["Age"] = titanic["Age"].fillna(titanic["Age"].median())
titanic["Embarked"].fillna('S', inplace = True)
titanic.isnull().sum()
```

Out[6]:

```
Survived      0
Pclass        0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
dtype: int64
```

Visualization

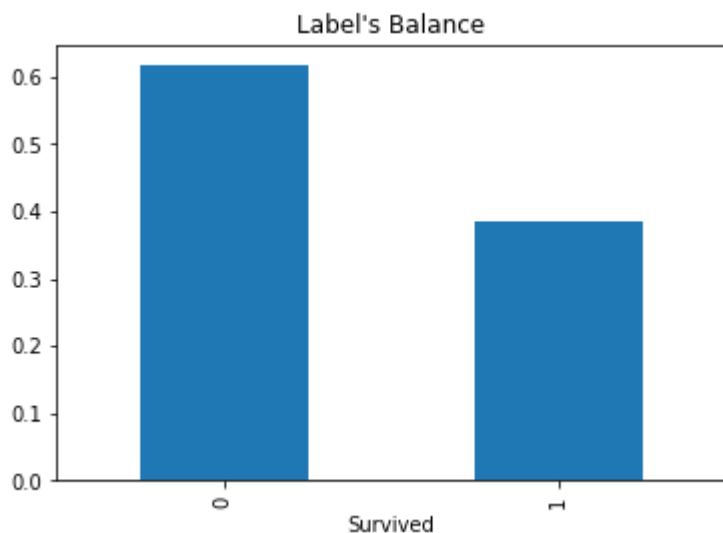
Entrée [7]:

```
%matplotlib inline
import matplotlib.pyplot as plt
print ('survival rate =', titanic.Survived.mean())
(titanic.groupby('Survived').size() / titanic.shape[0]).plot(kind = "bar", title = "
```

survival rate = 0.3838383838383838

Out[7]:

<AxesSubplot:title={'center':"Label's Balance"}, xlabel='Survived'>



Entrée [8]:


```
# make a function to plot survival against passenger attribute
def survival_rate(column, t):
    df = pd.DataFrame()
    df['total'] = titanic.groupby(column).size()
    df['survived'] = titanic.groupby(column).sum()['Survived']
    df['percentage'] = round(df['survived'] / df['total'] * 100, 2)
    print(df)

    df['survived'].plot(kind = t)
    df['total'].plot(kind = t, alpha = 0.5, title = "Survivors per " + str(column))
    plt.show()
```

Entrée [9]:

```
# Draw survival per Sex  
survival_rate("Sex", "barh")
```

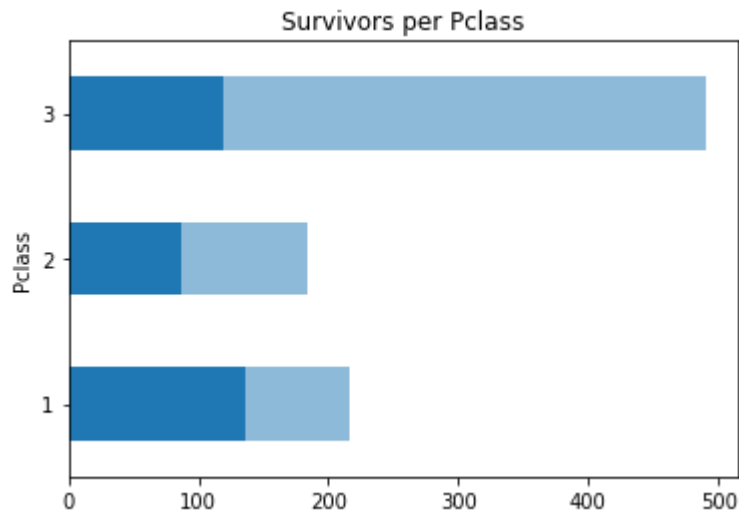
	total	survived	percentage
Sex			
female	314	233	74.20
male	577	109	18.89



Entrée [10]:

```
# Draw survival per Class
survival_rate("Pclass", "barh")
```


	total	survived	percentage
Pclass			
1	216	136	62.96
2	184	87	47.28
3	491	119	24.24



Entrée [11]:

```
# Graph survived per port of embarkation
survival_rate("Embarked", "bar")
```

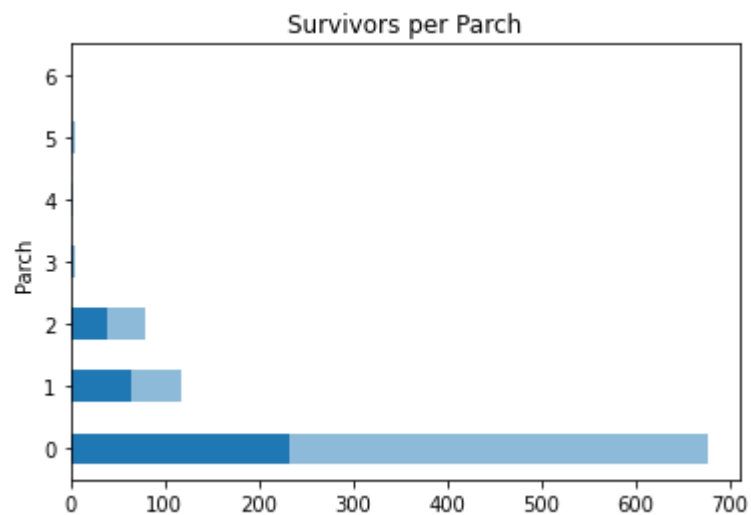
	total	survived	percentage
Embarked			
C	168	93	55.36
Q	77	30	38.96
S	646	219	33.90



Entrée [12]:

```
# Draw survived per Number of Parents/Children Aboard (Parch)
survival_rate("Parch", "barh")
```

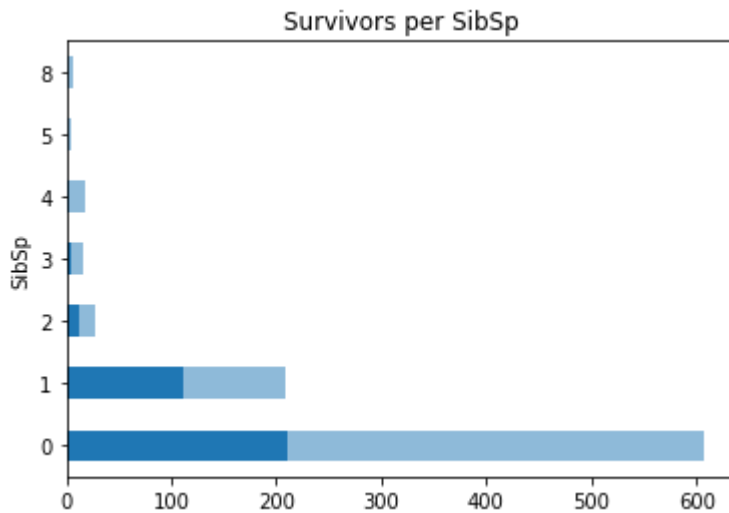
	total	survived	percentage
Parch			
0	678	233	34.37
1	118	65	55.08
2	80	40	50.00
3	5	3	60.00
4	4	0	0.00
5	5	1	20.00
6	1	0	0.00



Entrée [13]:

```
# Draw survived per Number of Siblings/Spouses Aboard (SibSp)
survival_rate("SibSp", "barh")
```

	total	survived	percentage
SibSp			
0	608	210	34.54
1	209	112	53.59
2	28	13	46.43
3	16	4	25.00
4	18	3	16.67
5	5	0	0.00
8	7	0	0.00



Model training

Some of the columns don't have predictive power, so let's specify which ones are included for prediction

Entrée [14]:

```
predictors = ["Pclass", "Sex", "Age", "SibSp", "Parch", "Fare", "Embarked"]
```

We need now to convert text columns in **predictors** to numerical ones

Entrée [15]:

```

for col in predictors: # Loop through all columns in predictors
    if titanic[col].dtype == 'object': # check if column's type is object (text)
        titanic[col] = pd.Categorical(titanic[col]).codes # convert text to numeric

titanic.head()

```

Out[15]:

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embark
PassengerId										
1	0	3	Braund, Mr. Owen Harris	1	22.0	1	0	A/5 21171	7.2500	
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	0	38.0	1	0	PC 17599	71.2833	
3	1	3	Heikkinen, Miss. Laina	0	26.0	0	0	STON/O2. 3101282	7.9250	
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	0	35.0	1	0	113803	53.1000	
5	0	3	Allen, Mr. William Henry	1	35.0	0	0	373450	8.0500	

Entrée [16]:

```

# Split the data into a training set and a testing set. Set: test_size=0.3, random_s
from sklearn.model_selection import train_test_split
# your code here

y = titanic["Survived"]
X = titanic[predictors]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_st

print ("train shape", X_train.shape, y_train.shape)
print ("test shape", X_test.shape, y_test.shape)

train shape (623, 7) (623,)
test shape (268, 7) (268,)

```

Entrée [17]:

```
# import LogisticRegression from: http://scikit-learn.org/stable/modules/generated/s

# your code here
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state = 1)
clf = clf.fit(X_train, y_train)

# your code here
train_score = clf.score(X_train, y_train)
test_score = clf.score(X_test, y_test)

print ('train accuracy =', train_score)
print ('test accuracy =', test_score)
```

```
train accuracy = 0.8073836276083467
test accuracy = 0.7723880597014925
```

Let's print the model's parameters

Entrée [18]:

```
coeff = pd.DataFrame()
coeff['Feature'] = X_train.columns
coeff['Coefficient Estimate'] = pd.Series(clf.coef_[0])
coeff.loc[len(coeff)] = ['Intercept', clf.intercept_[0]]
print(coeff)
```

	Feature	Coefficient Estimate
0	Pclass	-1.158693
1	Sex	-2.708761
2	Age	-0.040634
3	SibSp	-0.334012
4	Parch	0.071939
5	Fare	-0.000570
6	Embarked	-0.223307
7	Intercept	5.391545

We now need to predict class labels for the test set. We will also generate the class probabilities

Entrée [19]:

```
# predict class labels for the test set
y_pred = clf.predict(X_test)
print (y_pred)
```

```
[1 0 1 1 1 0 0 1 1 1 0 1 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 1 0 1 1 0 0 1 0
1 0
 0 1 0 1 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 0
0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 1 0 0 1 1 0 0 0 0
0 0
 0 0 0 0 0 0 0 1 1 1 0 0 0 1 1 1 1 0 0 0 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1
0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 1 1 0 0 0 1 1 1 0 1 0 0 0 1 0 1 1 0
0 1
 0 0 1 0 1 0 0 1 1 1 1 0 1 0 0 0 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 1 1 0
0 0
 0 0 0 0 1 0 1 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 1 1 0 1 0 0 0 1 1 1 0
1 0
 1 1 0 1 1 0 0 1 0]
```

Entrée [20]:

```
# generate class probabilities : http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.LogisticRegression.predict\_proba.html
y_probs = clf.predict_proba(X_test)
print (y_probs)
```

```
[[0.13931138 0.86068862]
 [0.91545699 0.08454301]
 [0.12666421 0.87333579]
 [0.36591842 0.63408158]
 [0.07784301 0.92215699]
 [0.89067951 0.10932049]
 [0.80772606 0.19227394]
 [0.11922783 0.88077217]
 [0.49862954 0.50137046]
 [0.44730682 0.55269318]
 [0.90553194 0.09446806]
 [0.39942792 0.60057208]
 [0.75883959 0.24116041]
 [0.7731844 0.2268156 ]
 [0.36586717 0.63413283]
 [0.51884413 0.48115587]
 [0.92555342 0.07444658]
 [0.92813017 0.07186983]
 [0.93745225 0.06254775]
 ...]
```

As you can see, the classifier outputs two probabilities for each row. It's predicting a 1 (Survived) any time the probability in the second column is greater than 0.5. Let's visualize it all together.

Entrée [21]:

```
pred = pd.DataFrame({
    "Survived_original": y_test,
    "Survived_predicted": y_pred,
    "Survived_proba": np.transpose(y_probs)[1]
})
pred["Comparison"] = pred.Survived_original == pred.Survived_predicted
pred.head()
```

Out[21]:

	Survived_original	Survived_predicted	Survived_proba	Comparison
PassengerId				
863	1	1	0.860689	True
224	0	0	0.084543	True
85	1	1	0.873336	True
681	0	1	0.634082	False
536	1	1	0.922157	True

Confusion matrix

Entrée [22]:

```
from sklearn import metrics
print (metrics.confusion_matrix(y_test, y_pred))
print (metrics.classification_report(y_test, y_pred))
```

```
[[129  24]
 [ 37  78]]
      precision    recall  f1-score   support

     0       0.78      0.84      0.81       153
     1       0.76      0.68      0.72       115

 accuracy          0.77
 macro avg          0.77      0.76      0.76       268
 weighted avg          0.77      0.77      0.77       268
```

As you can see, we can have the classification report for each class

K-Fold Cross Validation

Entrée [23]:

```
# import cross_validation from: http://scikit-learn.org/stable/modules/generated/sklearn.cross_validation
# your code here
from sklearn.model_selection import cross_val_score
clf = LogisticRegression(random_state = 1)
scores = cross_val_score(clf, titanic[predictors], titanic["Survived"], scoring = 'accuracy')
## see model
print(scores)
# Take the mean of the scores (because we have one for each fold)
print(scores.mean())
```

```
[0.7877095  0.78651685 0.78089888 0.76966292 0.82022472]
0.7890025735986442
```

When you are improving a model, you want to make sur that you are really doing it and not just being lucky. This is why it's good to work with cross validation instead of one train/test split.

Entrée []: